

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import itertools
```

## 1. Compare your model performance with different numbers of hidden nodes

```
In [3]: # Import dataset. Critical temperature unit: K
dataset = pd.read_csv('train.csv')      # Main feature dataset
formula = pd.read_csv('unique_m.csv')    # Formula of materials.
```

```
In [78]: x = dataset.iloc[:,0:dataset.shape[1]-1]
y = dataset["critical_temp"].values.reshape(x.shape[0], 1)
y_true = dataset["critical_temp"]
x.shape
```

```
Out[78]: (21263, 81)
```

Split data and standarization.

```
In [85]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

#split data into train and test set
Xtrain, Xtest, ytrain, ytest, ytrue_train, ytrue_test = train_test_split(np.array(x), np
```

```
In [80]: sc = StandardScaler()
sc.fit(Xtrain)
Xtrain = sc.transform(Xtrain)
Xtest = sc.transform(Xtest)

print(f"Shape of train set is {Xtrain.shape}")
print(f"Shape of test set is {Xtest.shape}")
print(f"Shape of train label is {ytrain.shape}")
print(f"Shape of test labels is {ytest.shape}")
```

```
Shape of train set is (17010, 81)
Shape of test set is (4253, 81)
Shape of train label is (17010, 1)
Shape of test labels is (4253, 1)
```

```
In [269... class NeuralNet():
    '''
    A two layer neural network
    '''

    def __init__(self, layers=[13,8,1], learning_rate=0.001, iterations=100):
        self.params = {}
        self.learning_rate = learning_rate
```

```

self.iterations = iterations
self.loss = []
self.sample_size = None
self.layers = layers
self.X = None
self.y = None

def init_weights(self):
    """
    Initialize the weights from a random normal distribution
    """
    np.random.seed(1) # Seed the random number generator
    for i in range(len(self.layers)-1):
        ws = "W"+str(i+1)
        bs = "b"+str(i+1)
        self.params[ws] = np.random.randn(self.layers[i], self.layers[i+1])
        self.params[bs] = np.random.randn(self.layers[i+1],)
    # self.params["W1"] = np.random.randn(self.layers[0], self.layers[1])
    # self.params['b1'] = np.random.randn(self.layers[1],)
    # self.params['W2'] = np.random.randn(self.layers[1],self.layers[2])
    # self.params['b2'] = np.random.randn(self.layers[2],)

def relu(self, Z):
    """
    The ReLu activation function is to performs a threshold
    operation to each input element where values less
    than zero are set to zero.
    """
    return np.maximum(0, Z)

def dRelu(self, x):
    x[x<=0] = 0
    x[x>0] = 1
    return x

def eta(self, x):
    ETA = 0.0000000001
    return np.maximum(x, ETA)

def tanh(self, Z):
    # t=(np.exp(Z)-np.exp(-Z))/(np.exp(Z)+np.exp(-Z))
    t = np.tanh(Z)
    return t

def dtanh(self, Z):
    t=(np.exp(Z)-np.exp(-Z))/(np.exp(Z)+np.exp(-Z))
    dt=1-t**2
    return dt

def linear(self, Z):
    return Z

def dlinear(self, Z):
    return 1

def sigmoid(self, Z):
    """
    The sigmoid function takes in real numbers in any range and
    squashes it to a real-valued output between 0 and 1.
    """
    return 1/(1+np.exp(-Z))

def entropy_loss(self, y, yhat):
    nsample = len(y)
    yhat_inv = 1.0 - yhat
    y_inv = 1.0 - y

```

```

yhat = self.eta(yhat) ## clips value to avoid NaNs in log
yhat_inv = self.eta(yhat_inv)
loss = -1/nsample * (np.sum(np.multiply(np.log(yhat), y) + np.multiply((y_inv),
return loss

def forward_propagation(self):
'''
Performs the forward propagation
'''
# print(f"self params shape:{self.params['W1']}")
Z1 = self.X.dot(self.params['W1']) + self.params['b1']
A1 = self.relu(Z1)
Z2 = A1.dot(self.params['W2']) + self.params['b2']
yhat = self.sigmoid(Z2)
loss = self.entropy_loss(self.y, yhat)

# save calculated parameters
self.params['Z1'] = Z1
self.params['Z2'] = Z2
self.params['A1'] = A1

return yhat, loss

def forward_propagation2(self):
'''
Performs the forward propagation
'''
# print(f"self params shape:{self.params['W1']}")
Z1 = self.X.dot(self.params['W1']) + self.params['b1']
A1 = self.tanh(Z1)
Z2 = A1.dot(self.params['W2']) + self.params['b2']
yhat = self.linear(Z2)
loss = self.entropy_loss(self.y, yhat)

# save calculated parameters
self.params['Z1'] = Z1
self.params['Z2'] = Z2
self.params['A1'] = A1

return yhat, loss

def back_propagation(self, yhat):
'''
Computes the derivatives and update weights and bias according.
'''
y_inv = 1 - self.y
yhat_inv = 1 - yhat

dl_wrt_yhat = np.divide(y_inv, self.eta(yhat_inv)) - np.divide(self.y, self.eta(
dl_wrt_sig = yhat * (yhat_inv)
dl_wrt_z2 = dl_wrt_yhat * dl_wrt_sig

dl_wrt_A1 = dl_wrt_z2.dot(self.params['W2']).T)
dl_wrt_w2 = self.params['A1'].T.dot(dl_wrt_z2)
dl_wrt_b2 = np.sum(dl_wrt_z2, axis=0, keepdims=True)

dl_wrt_z1 = dl_wrt_A1 * self.dRelu(self.params['Z1'])
dl_wrt_w1 = self.X.T.dot(dl_wrt_z1)
dl_wrt_b1 = np.sum(dl_wrt_z1, axis=0, keepdims=True)

#update the weights and bias
self.params['W1'] = self.params['W1'] - self.learning_rate * dl_wrt_w1
self.params['W2'] = self.params['W2'] - self.learning_rate * dl_wrt_w2
self.params['b1'] = self.params['b1'] - self.learning_rate * dl_wrt_b1
self.params['b2'] = self.params['b2'] - self.learning_rate * dl_wrt_b2

```

```

def back_propagation2(self, yhat):
    """
    Computes the derivatives and update weights and bias according.
    """
    y_inv = 1 - self.y
    yhat_inv = 1 - yhat

    dl_wrt_yhat = np.divide(y_inv, self.eta(yhat_inv)) - np.divide(self.y, self.eta(
    dl_wrt_sig = yhat * (yhat_inv)
    dl_wrt_z2 = dl_wrt_yhat * 1

    dl_wrt_A1 = dl_wrt_z2.dot(self.params['W2']).T
    dl_wrt_w2 = self.params['A1'].T.dot(dl_wrt_z2)
    dl_wrt_b2 = np.sum(dl_wrt_z2, axis=0, keepdims=True)

    dl_wrt_z1 = dl_wrt_A1 * self.dtanh(self.params['Z1'])
    dl_wrt_w1 = self.X.T.dot(dl_wrt_z1)
    dl_wrt_b1 = np.sum(dl_wrt_z1, axis=0, keepdims=True)

    #update the weights and bias
    self.params['W1'] = self.params['W1'] - self.learning_rate * dl_wrt_w1
    self.params['W2'] = self.params['W2'] - self.learning_rate * dl_wrt_w2
    self.params['b1'] = self.params['b1'] - self.learning_rate * dl_wrt_b1
    self.params['b2'] = self.params['b2'] - self.learning_rate * dl_wrt_b2

def fit(self, X, y):
    """
    Trains the neural network using the specified data and labels
    """
    self.X = X
    self.y = y
    self.init_weights() #initialize weights and bias

    for i in range(self.iterations):
        yhat, loss = self.forward_propagation()
        self.back_propagation(yhat)
        self.loss.append(loss)

def fit2(self, X, y):
    """
    Trains the neural network using the specified data and labels
    """
    self.X = X
    self.y = y
    self.init_weights() #initialize weights and bias

    for i in range(self.iterations):
        yhat, loss = self.forward_propagation2()
        # self.back_propagation2(yhat)
        self.loss.append(loss)

def predict(self, X):
    """
    Predicts on a test data
    """
    Z1 = X.dot(self.params['W1']) + self.params['b1']
    A1 = self.relu(Z1)
    Z2 = A1.dot(self.params['W2']) + self.params['b2']
    pred = self.sigmoid(Z2)
    return np.round(pred)

def predict2(self, X):
    """
    Predicts on a test data

```

```

'''
Z1 = X.dot(self.params['W1']) + self.params['b1']
A1 = self.tanh(Z1)
Z2 = A1.dot(self.params['W2']) + self.params['b2']
pred = self.linear(Z2)
return np.round(pred)

def acc(self, y, yhat):
'''
Calculates the accuracy between the predicted values and actual
'''
acc = int(sum(y == yhat) / len(y) * 100)
return acc

def plot_loss(self):
'''
Plots the loss curve
'''
plt.plot(self.loss)
plt.xlabel("Iteration")
plt.ylabel("logloss")
plt.title("Loss curve for training")
plt.show()

```

```

In [281... nn1 = NeuralNet(layers=[Xtrain.shape[1], 10, 1],
learning_rate=0.001, iterations=100) # create the NN model
nn1.fit(Xtrain, ytrain) #train the model

```

C:\Users\youyu\AppData\Local\Temp\ipykernel\_31640\2579518989.py:69: RuntimeWarning: overflow encountered in exp  
return 1/(1+np.exp(-Z))

```

In [111... train_pred1 = nn1.predict(Xtrain)
test_pred1 = nn1.predict(Xtest)

print("Train accuracy is {}".format(nn1.acc(ytrain, train_pred1)))
print("Test accuracy is {}".format(nn1.acc(ytest, test_pred1)))

```

Train accuracy is 0  
Test accuracy is 0

C:\Users\youyu\AppData\Local\Temp\ipykernel\_31640\1925195097.py:61: RuntimeWarning: overflow encountered in exp  
return 1/(1+np.exp(-Z))

```

In [92]: cost_train1 = (ytrain-train_pred1)**2
prediction_train1 = pd.DataFrame({"Y": ytrue_train,
"Y_predict": train_pred1.reshape(train_pred1.shape[0]),
"cost":cost_train1.reshape(cost_train1.shape[0])})
prediction_train1

```

Out[92]:

	Y	Y_predict	cost
11065	0.50	1.0	0.2500
5689	40.00	1.0	1521.0000
2176	84.00	1.0	6889.0000
9595	32.20	1.0	973.4400
9034	78.10	1.0	5944.4100
...	...	...	...
1099	92.50	1.0	8372.2500

<b>18898</b>	17.10	1.0	259.2100
<b>11798</b>	0.21	1.0	0.6241
<b>6637</b>	88.40	1.0	7638.7600
<b>2575</b>	56.00	1.0	3025.0000

17010 rows × 3 columns


```
In [282... rmse_train1 = np.sqrt(sum(cost_train1)/len(cost_train1))
r2_train1= 1-sum(cost_train1)/sum((np.array(train_pred1.reshape(train_pred1.shape[0]))-s

print(f'RMSE of NN on training dataset = {round(rmse_train1[0],3)}')
print(f'R2 score = {round(r2_train1[0],3)}')
```


RMSE of NN on training dataset = 47.995  
R2 score = -inf

C:\Users\youyu\AppData\Local\Temp\ipykernel\_31640\2353467826.py:2: RuntimeWarning: divide by zero encountered in true\_divide  
r2\_train1= 1-sum(cost\_train1)/sum((np.array(train\_pred1.reshape(train\_pred1.shape[0]))-sum(np.array(train\_pred1.reshape(train\_pred1.shape[0])))/len(np.array(train\_pred1.reshape(train\_pred1.shape[0]))))\*2)

```
In [94]: plt.scatter(prediction_train1['Y'], prediction_train1['Y_predict'], label='True')
plt.plot(prediction_train1['Y'], prediction_train1['Y'], label='Perfect prediction when
plt.legend(loc="upper left")
plt.title('Compare Train and Prediction')
plt.xlabel('Target Value: True Critical Temperature (k)')
plt.ylabel('Predicted Critical Temperature')
plt.show()
```



```
In [102... nn2 = NeuralNet(layers=[Xtrain.shape[1], 5, 1],
learning_rate=0.001, iterations=100) # create the NN model
nn2.fit(Xtrain, ytrain) #train the model
nn2.plot_loss()
```



```
In [118... train_pred2 = nn2.predict(Xtrain)
test_pred2 = nn2.predict(Xtest)

print("Train accuracy is {}".format(nn2.acc(ytrain, train_pred2)))
print("Test accuracy is {}".format(nn2.acc(ytest, test_pred2)))
```

```
Train accuracy is 0
Test accuracy is 0
```

```
In [293... cost_train2 = (ytrain-train_pred2)**2
rmse_train2 = np.sqrt(sum(cost_train2)/len(cost_train2))
r2_train2= 1-sum(cost_train2)/sum((np.array(train_pred2.reshape(train_pred2.shape[0]))-s

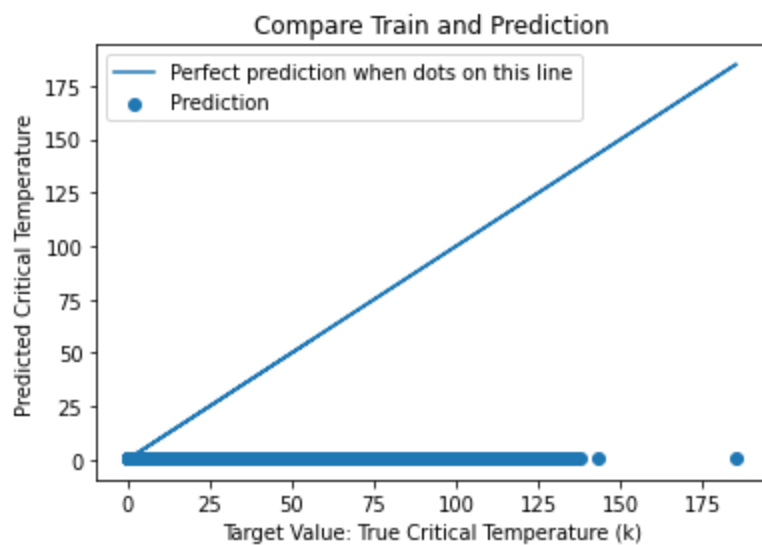
print(f'RMSE of NN on training dataset = {round(rmse_train2[0],3)}')
print(f'R2 score = {round(r2_train2[0],3)}')
```

```
RMSE of NN on training dataset = 47.995
R2 score = -inf
```

```
C:\Users\youyu\AppData\Local\Temp\ipykernel_31640\1011046090.py:3: RuntimeWarning: divid
e by zero encountered in true_divide
  r2_train2= 1-sum(cost_train2)/sum((np.array(train_pred2.reshape(train_pred2.shape[0]))
-sum(np.array(train_pred2.reshape(train_pred2.shape[0]))/len(np.array(train_pred2)))**
2)
```

```
In [276... cost_train2 = (ytrain-train_pred2)**2
prediction_train2 = pd.DataFrame({"Y": ytrue_train,
    "Y_predict": train_pred2.reshape(train_pred2.shape[0]),
    "cost":cost_train2.reshape(cost_train2.shape[0])})

plt.scatter(prediction_train2['Y'], prediction_train2['Y_predict'], label='Prediction')
plt.plot(prediction_train2['Y'], prediction_train2['Y'], label='Perfect prediction when
plt.legend(loc="upper left")
plt.title('Compare Train and Prediction')
plt.xlabel('Target Value: True Critical Temperature (k)')
plt.ylabel('Predicted Critical Temperature')
plt.show()
```



The NNs above use 1 hidden layer with 10 and 5 nodes between. The first activation function is ReLU and the second one is sigmoid function. With the current project data cleaning process and modeling, the outcome accuracy is 0. The activation functions did not work in this example. The prediction accuracy and the RMSE, R2 score did not fit the data at all. All the prediction values are 0, which is not the prediction value should be.

At this time, we can conclude that the neural network with ReLU and sigmoid functions are not ideal for critical temperature prediction in this case.

## 2. Show the model performance with two different activation functions

This part will use tanh and linear activation function to fit and predict.

```
In [270...] nn3 = NeuralNet(layers=[Xtrain.shape[1], 5, 1],
    learning_rate=0.001, iterations=100) # create the NN model
nn3.fit2(Xtrain, ytrain) #train the model
```

```
In [271...] train_pred3 = nn3.predict2(Xtrain)
test_pred3 = nn3.predict2(Xtest)

print("Train accuracy is {}".format(nn1.acc(ytrain, train_pred3)))
print("Test accuracy is {}".format(nn1.acc(ytest, test_pred3)))
```

Train accuracy is 0  
Test accuracy is 0

```
In [273...] cost_train3 = (ytrain-train_pred3)**2
rmse_train3 = np.sqrt(sum(cost_train3)/len(cost_train3))
r2_train3= 1-sum(cost_train3)/sum((np.array(train_pred3.reshape(train_pred3.shape[0]))-s

print(f'RMSE of tanh&linear model on training dataset = {round(rmse_train3[0],3)}')
print(f'R2 score = {round(r2_train3[0],3)}')
```


RMSE of tanh&linear model on training dataset = 45.302  
R2 score = -8648.547

```
In [278...] prediction_train3 = pd.DataFrame({"Y": ytrue_train,
    "Y_predict": train_pred3.reshape(train_pred3.shape[0]),
    "cost":cost_train3.reshape(cost_train3.shape[0])})

plt.scatter(prediction_train3['Y'], prediction_train3['Y_predict'], label='Prediction')
plt.plot(prediction_train3['Y'], prediction_train3['Y'], label='Perfect prediction when
plt.legend(loc="upper left")
plt.title('Compare Train and Prediction')
```



```
plt.xlabel('Target Value: True Critical Temperature (k)')
plt.ylabel('Predicted Critical Temperature')
plt.show()
```



```
In [283... prediction_train3
```

```
Out[283]:
```

	Y	Y_predict	cost
<b>11065</b>	0.50	5.0	20.2500
<b>5689</b>	40.00	5.0	1225.0000
<b>2176</b>	84.00	5.0	6241.0000
<b>9595</b>	32.20	5.0	739.8400
<b>9034</b>	78.10	5.0	5343.6100
...	...	...	...
<b>1099</b>	92.50	5.0	7656.2500
<b>18898</b>	17.10	5.0	146.4100
<b>11798</b>	0.21	5.0	22.9441
<b>6637</b>	88.40	5.0	6955.5600
<b>2575</b>	56.00	5.0	2601.0000

17010 rows × 3 columns

The prediction does not look good. The prediction maximum is obviously lower than most of the true values. R2 score indicated that the prediction did not correspond to the trend of true values. The tanh activation function would predict all the observations into similar values. The back\_propagation part does not work well on the process.

Overall, this model did not able to find the pattern of the dataset.

### 3. Show your output with Stochastic gradient and gradient optimization methods.

```
In [161... # Define the model parameters
learning_rate = 5e-11
num_epochs = 1000

# Initialize the model coefficients
```

```

theta = np.zeros(Xtrain.shape[1])

# Define the SGD model
def model(X, theta):
    return X @ theta

# Define the loss function
def loss(X, y, theta_0, theta_1):
    return 1 / len(X) * sum((model(X, theta_0, theta_1) - y) ** 2)


# Define the gradient descent step
def step(X, y, theta, learning_rate):
    gradients = 2 / len(X) * X.T @ (model(X, theta) - y)
    theta -= learning_rate * gradients
    return theta

# Perform gradient descent
for epoch in range(num_epochs):
    theta = step(Xtrain, ytrain.reshape(ytrain.shape[0]), theta, learning_rate)

# Plot the data and the model predictions
plt.scatter(ytrain.reshape(ytrain.shape[0]),
            ytrain.reshape(ytrain.shape[0]),
            color='blue')
plt.scatter(ytrain.reshape(ytrain.shape[0]), model(Xtrain, theta), color='red')
plt.xlabel('Ytrain')
plt.ylabel('Y')
plt.show()

# Print coefs
print('theta:', theta)

```



```

theta: [ 4.28373448e-06  4.58476446e-05  2.55660243e-05  3.46612158e-05
 1.71958529e-05  1.14855344e-06  1.18682683e-06  9.83444301e-05
-1.76314887e-06  3.06157297e-05  2.68670590e-05  4.85161081e-04
 6.39253830e-04  4.46486677e-04  5.87097899e-04  1.21698200e-06
 8.03429365e-07  6.90900047e-04  4.24039681e-04  2.34856291e-04
 2.68552870e-04  1.10509265e-04  7.68704186e-05  9.18817539e-05
 5.71719086e-05  1.18826832e-06  1.16217412e-06  1.71316164e-04
 1.16104343e-05  5.52631860e-05  5.93801947e-05  9.21426742e-04
 2.46493217e-04 -1.61890761e-03 -1.81532933e-03  1.04654406e-06
 7.74071086e-07  3.03312402e-03  4.12477654e-05  5.64908566e-04
 1.00955934e-03  2.86184189e-05  5.72653985e-05  4.11645611e-06
 3.05874224e-05  9.15476827e-07  5.50913574e-07  9.28117990e-05
 4.24727152e-05  3.69192053e-05  3.59176700e-05 -3.73149518e-07
-2.50373399e-06 -2.41806502e-06 -4.32845394e-06  1.12292261e-06
 9.75433869e-07  3.58727204e-06 -1.68232058e-06  1.85843930e-07
 2.12395306e-07  9.25989528e-05  1.01166377e-04 -5.42177960e-06
-1.09565621e-05  4.86113518e-07  1.96649657e-07  3.98529969e-04]

```

```

9.29446099e-05 1.49446657e-04 1.67112783e-04 5.82112993e-07
3.32542395e-07 6.29754316e-07 3.92258184e-07 1.27107164e-06
1.07234221e-06 2.04878585e-07 4.95048011e-08 3.91097713e-08
-1.11915502e-07]

```

```

In [189... train_pred4 = model(Xtrain, theta)
test_pred4 = model(Xtest, theta)

cost_train4 = (ytrain.reshape(ytrain.shape[0]) - train_pred4)**2
rmse_train4 = np.sqrt(sum(cost_train4)/len(cost_train4))
r2_train4 = 1 - sum(cost_train4)/sum((np.array(train_pred4) - sum(np.array(train_pred4))/len(

print(f'RMSE of stochastic gradient model on training dataset = {round(rmse_train4,3)}')
print(f'R2 score = {round(r2_train4,3)}')

RMSE of stochastic gradient model on training dataset = 30.897
R2 score = -1.209

```

The RMSE means that the cost of the prediction is not high, but R2 score indicated that the prediction values are not corresponding to the true patterns. The distribution of predicted values indicated that the model caught the target values' feature, but failed to find the relationships between X and Y. This could occurred especially when X has too many columns while the most correlated features were not emphasized.